

Stochastic Hyperparameter Optimization through Hypernetworks

Anonymous Author(s)

Affiliation

Address

email

Abstract

Machine learning models are often tuned by nesting optimization of model weights inside the optimization of hyperparameters. We give a method to collapse this nested optimization into joint stochastic optimization of weights and hyperparameters. Our process trains a neural network to output approximately optimal weights as a function of hyperparameters. We show that our technique converges to locally optimal weights and hyperparameters for sufficiently large hypernetworks. We compare this method to standard hyperparameter optimization strategies and demonstrate its effectiveness for tuning thousands of hyperparameters.

1 Introduction

Model selection and hyperparameter tuning is a significant bottleneck in designing predictive models. Hyperparameter optimization is a nested optimization: The inner optimization finds model parameters w which minimize the training loss $\mathcal{L}_{\text{Train}}$ given hyperparameters λ . The outer optimization chooses λ to reduce a validation loss $\mathcal{L}_{\text{Valid.}}$:

$$\operatorname{argmin}_{\lambda} \mathcal{L}_{\text{Valid.}} \left(\operatorname{argmin}_{w} \mathcal{L}_{\text{Train}}(w, \lambda) \right) \quad (1)$$

Standard practice in machine learning solves (1) by gradient-free optimization of hyperparameters, such as grid search or random search. Each set of hyperparameters is evaluated by re-initializing weights and training the model to completion. Retraining a model from scratch is wasteful if the hyperparameters change by a small amount. Some approaches, such as Hyperband (Li et al., 2016) and freeze-thaw Bayesian optimization (Swersky et al., 2014), can pause model training and not waste this effort. However, these methods often scale poorly beyond 10 to 20 dimensions.

How can we avoid re-training from scratch each time? Note that the optimal parameters w are a deterministic function of the hyperparameters λ :

$$w^*(\lambda) = \operatorname{argmin}_{w} \mathcal{L}_{\text{Train}}(w, \lambda) \quad (2)$$




Figure 1: *Left*: A typical computational graph for cross-validation, where α are the optimizer parameters, and λ are training loss hyperparameters. It is expensive to differentiate through the entire training procedure. *Right*: The proposed computational graph with our changes in red, where ϕ are the hypernetwork parameters. We can cheaply differentiate through the hypernetwork to optimize the validation loss $\mathcal{L}_{\text{Valid.}}$ with respect to hyperparameters λ . We use x , t , and y to refer to a data point, its label, and a prediction respectively.




Figure 2: A visualization of exact (blue) and approximate (red) optimal weights as a function of hyperparameters. The approximately optimal weights w_{ϕ^*} are output by a linear model fit at $\hat{\lambda}$. The true optimal hyperparameter is λ^* , while the hyperparameter estimated using approximately optimal weights is nearby at λ_{ϕ^*} .

32 We propose to *learn this function*. Specifically, we train a neural network that takes hyperparameters
 33 as input, and outputs an approximately optimal set of weights.

34 This formulation provides two major benefits: First, we can train the hypernetwork to convergence
 35 using stochastic gradient descent (SGD) without training any particular model to completion. Sec-
 36 ond, differentiating through the hypernetwork allows us to optimize hyperparameters with stochastic
 37 gradient-based optimization.

38 2 Training a network to output optimal weights

39 How can we teach a *hypernetwork* (Ha et al., 2016) to output approximately optimal weights to
 40 another neural network? The basic idea is that at each iteration, we ask a hypernetwork to output a
 41 set of weights given some hyperparameters: $w = w_{\phi}(\lambda)$. Instead of updating the weights w using
 42 the training loss gradient $\partial \mathcal{L}_{\text{Train}}(w)/\partial w$, we update the hypernetwork weights ϕ using the chain
 43 rule: $\partial \mathcal{L}_{\text{Train}}(w_{\phi})/\partial w_{\phi} \partial w_{\phi}/\partial \phi$. This formulation allows us to optimize the hyperparameters λ with
 44 the validation loss gradient $\partial \mathcal{L}_{\text{Valid.}}(w_{\phi}(\lambda))/\partial w_{\phi}(\lambda) \partial w_{\phi}(\lambda)/\partial \lambda$. We call this method *hyper-training* and
 45 contrast it with standard training methods.

46 We call the function $w^*(\lambda)$ that outputs optimal weights for hyperparameters a *best-response func-*
 47 *tion*. At convergence, we want our hypernetwork $w_{\phi}(\lambda)$ to match the best-response function closely.

48 Our method is closely related to the concurrent work of Brock et al. (2017), whose SMASH al-
 49 gorithm also approximates the optimal weights as a function of model architectures, to perform a
 50 gradient-free search over discrete model structures. Their work focuses on efficiently estimating
 51 the performance of a variety of model architectures, while we focus on efficiently exploring contin-
 52 uous spaces of models. We further extend this idea by formulating an algorithm to optimize the
 53 hypernetwork and hyperparameters jointly. Joint optimization of parameters and hyperparameters
 54 addresses one of the main weaknesses of SMASH, which is that the hypernetwork must be very
 55 large to learn approximately optimal weights for many different settings. During joint optimization,
 56 the hypernetwork need only model approximately optimal weights for the neighborhood around the
 57 current hyperparameters, allowing us to use even linear hypernetworks.

58 2.1 Advantages of hypernetwork-based optimization

59 Hyper-training is a method to learn a mapping from hyperparameters to validation loss which is
 60 differentiable and cheap to evaluate. We can compare hyper-training to other model-based hyperpa-
 61 rameter schemes. Bayesian optimization (e.g., Lizotte (2008); Snoek et al. (2012)) builds a model
 62 of the validation loss as a function of hyperparameters, usually using a Gaussian process (e.g., Ras-
 63 mussen & Williams (2006)) to track uncertainty. This approach has several disadvantages compared
 64 to hyper-training.

First, obtaining data for standard Bayesian optimization requires optimizing models from initialization for each set of hyperparameters. In contrast, hyper-training never needs to optimize any one model fully, avoiding difficult choices such as how many models to train and for how long.

Second, standard Bayesian optimization treats the validation loss as a black-box function: $\mathcal{L}_{\text{Valid}}(\lambda) = f(\lambda)$. In contrast, hyper-training takes advantage of the fact that the validation loss is a known, differentiable function: $\mathcal{L}_{\text{Valid}}(\lambda) = \mathcal{L}_{\text{Valid}}(\mathbf{w}_\phi(\lambda))$. This information removes the need to learn a model of the validation loss. This function can also be evaluated stochastically by sampling points from the validation set.

Hyper-training has a benefit of learning hyperparameter to optimized weight mapping, which is substituted into the validation loss. This often has a better inductive bias for learning hyperparameter to validation loss than directly learning the loss. Also, the hypernetwork learns continuous best-responses, which may be a beneficial prior for finding weights by enforcing stability.

2.2 Limitations of hypernetwork-based optimization

We can apply this method to unconstrained continuous bi-level optimization problems with an inner loss function with inner parameters, and an outer loss function with outer parameters. What sort of parameters can be optimized by our approach? Hyperparameters typically fall into two broad categories: 1) Optimization hyperparameters, such as learning rates and initialization, which affect the choice of locally optimal point converged to, and 2) regularization or model architecture parameters which change the set of locally optimal points. Hyper-training *does not have inner optimization parameters* because there is no internal training loop, so we can not optimize these. However, we must still choose optimization parameters for the fused optimization loop. In principle, hyper-training can handle discrete hyperparameters, but does not offer particular advantages for optimization over continuous hyperparameters.

Another limitation is that our approach only proposes making local changes to the hyperparameters, and does not do uncertainty-based exploration. Uncertainty can be incorporated into the hypernetwork by using stochastic variational inference as in [Blundell et al. \(2015\)](#), and we leave this for future work. Finally, it is not obvious how to choose the training distribution of hyperparameters $p(\lambda)$. If we do not sample a sufficient range of hyperparameters, the implicit estimated gradient of the validation loss w.r.t. the hyperparameters may be inaccurate. We discuss several approaches to this problem in section 2.4.

A clear difficulty of this approach is that hypernetworks can require several times as many parameters as the original model. For example, training a fully-connected hypernetwork with 1 hidden layer of H units to output D parameters requires at least $D \times H$ hypernetwork parameters. To address this problem, in section 2.4, we propose an algorithm that only trains a linear model mapping hyperparameters to model weights.

2.3 Asymptotic convergence properties

Algorithm 2 trains a hypernetwork using SGD, drawing hyperparameters from a fixed distribution $p(\lambda)$. This section proves that Algorithm 2 converges to a local best-response under mild assumptions. In particular, we show that, for a sufficiently large hypernetwork, the choice of $p(\lambda)$ does not matter as long as it has sufficient support. Notation as if \mathbf{w}_ϕ has a unique solution for ϕ or \mathbf{w} is used for simplicity, but is not true in general.

Theorem 2.1. *Sufficiently powerful hypernetworks can learn continuous best-response functions, which minimizes the expected loss for all hyperparameter distributions with convex support.*

There exists ϕ^* , such that for all $\lambda \in \text{support}(p(\lambda))$,

$$\mathcal{L}_{\text{Train}}(\mathbf{w}_{\phi^*}(\lambda), \lambda) = \min_{\mathbf{w}} \mathcal{L}_{\text{Train}}(\mathbf{w}, \lambda) \quad \text{and} \quad \phi^* = \underset{\phi}{\operatorname{argmin}} \mathbb{E}_{p(\lambda')} \left[\mathcal{L}_{\text{Train}}(\mathbf{w}_\phi(\lambda'), \lambda') \right]$$

Proof. If \mathbf{w}_ϕ is a universal approximator ([Hornik, 1991](#)) and the best-response is continuous in λ (which allows approximation by \mathbf{w}_ϕ), then there exists optimal hypernetwork parameters ϕ^* such that for all hyperparameters λ , $\mathbf{w}_{\phi^*}(\lambda) = \operatorname{argmin}_{\mathbf{w}} \mathcal{L}_{\text{Train}}(\mathbf{w}, \lambda)$. Thus, $\mathcal{L}_{\text{Train}}(\mathbf{w}_{\phi^*}(\lambda), \lambda) = \min_{\mathbf{w}} \mathcal{L}_{\text{Train}}(\mathbf{w}, \lambda)$. In other words, universal approximator hypernetworks can learn continuous best-responses.

Algorithm 1 Standard cross-validation with stochastic optimization	Algorithm 2 Optimization of hypernetwork, then hyperparameters	Algorithm 3 Joint optimization of hypernetwork and hyperparameters
for $i = 1, \dots, T_{\text{outer}}$ do initialize w $\lambda = \Lambda(\lambda^{(1:i)}, \mathcal{L}_{\text{Valid.}}(w^{(1:i)}))$ loop $\mathbf{x} \sim \text{Training data}$ $w \leftarrow \alpha \nabla_w \mathcal{L}_{\text{Train}}(w, \lambda, \mathbf{x})$ $\lambda^i, w^i = \lambda, w$ $i = \underset{i}{\operatorname{argmin}} \mathcal{L}_{\text{Valid.}}(w^{(i)}, \mathbf{x})$ Return $\lambda^{(i)}, w^{(i)}$	initialize ϕ initialize $\hat{\lambda}$ loop $\mathbf{x} \sim \text{Training data}, \lambda \sim p(\lambda)$ $\phi \leftarrow \alpha \nabla_{\phi} \mathcal{L}_{\text{Train}}(w_{\phi}(\lambda), \lambda, \mathbf{x})$ loop $\mathbf{x} \sim \text{Validation data}$ $\hat{\lambda} \leftarrow \beta \nabla_{\hat{\lambda}} \mathcal{L}_{\text{Valid.}}(w_{\phi}(\hat{\lambda}), \mathbf{x})$ Return $\hat{\lambda}, w_{\phi}(\hat{\lambda})$	initialize ϕ initialize $\hat{\lambda}$ loop $\mathbf{x} \sim \text{Training data}, \lambda \sim p(\lambda \hat{\lambda})$ $\phi \leftarrow \alpha \nabla_{\phi} \mathcal{L}_{\text{Train}}(w_{\phi}(\lambda), \lambda, \mathbf{x})$ $\mathbf{x} \sim \text{Validation data}$ $\hat{\lambda} \leftarrow \beta \nabla_{\hat{\lambda}} \mathcal{L}_{\text{Valid.}}(w_{\phi}(\hat{\lambda}), \mathbf{x})$ Return $\hat{\lambda}, w_{\phi}(\hat{\lambda})$

A comparison of standard hyperparameter optimization, our first algorithm, and our joint algorithm. Here, Λ refers to a generic hyperparameter optimization. Instead of updating weights w using the loss gradient $\partial \mathcal{L}(w)/\partial w$, we update hypernetwork weights ϕ and hyperparameters λ using the chain rule: $\partial \mathcal{L}_{\text{Train}}(w_{\phi})/\partial w_{\phi} \partial w_{\phi}/\partial \phi$ or $\partial \mathcal{L}_{\text{Valid.}}(w_{\phi}(\lambda))/\partial w_{\phi}(\lambda) \partial w_{\phi}(\lambda)/\partial \lambda$ respectively. This allows our method to use gradient-based hyperparameter optimization.

113 Substituting ϕ^* into the training loss gives $\mathbb{E}_{p(\lambda)}[\mathcal{L}_{\text{Train}}(w_{\phi^*}(\lambda), \lambda)] = \mathbb{E}_{p(\lambda)}[\min_{\phi} \mathcal{L}_{\text{Train}}(w_{\phi}(\lambda), \lambda)]$.
 114 By Jensen’s inequality, $\min_{\phi} \mathbb{E}_{p(\lambda)}[\mathcal{L}_{\text{Train}}(w_{\phi}(\lambda), \lambda)] \geq \mathbb{E}_{p(\lambda)}[\min_{\phi} \mathcal{L}_{\text{Train}}(w_{\phi}(\lambda), \lambda)]$. To satisfy
 115 Jensen’s requirements, we have \min_{ϕ} as our convex function on the convex vector space of
 116 functions $\{\mathcal{L}_{\text{Train}}(w_{\phi}(\lambda), \lambda) \text{ for } \lambda \in \text{support}(p(\lambda))\}$. To guarantee convexity of the vector space
 117 we require that $\text{support}(p(\lambda))$ is convex and $\mathcal{L}_{\text{Train}}(w, \lambda) = \mathbb{E}_{\mathbf{x} \sim \text{Train}}[\mathcal{L}_{\text{Pred}}(\mathbf{x}, w)] + \mathcal{L}_{\text{Reg}}(w, \lambda)$
 118 with $\mathcal{L}_{\text{Reg}}(w, \lambda) = \lambda \cdot \mathcal{L}(w)$. Thus, $\phi^* = \operatorname{argmin}_{\phi} \mathbb{E}_{p(\lambda)}[\mathcal{L}_{\text{Train}}(w_{\phi}(\lambda), \lambda)]$. In other words, if the
 119 hypernetwork learns the best-response it will simultaneously minimize the loss for every point in
 120 $\text{support}(p(\lambda))$. \square

121 Thus, having a universal approximator and a continuous best-response implies for all $\lambda \in$
 122 $\text{support}(p(\lambda))$, $\mathcal{L}_{\text{Valid.}}(w_{\phi^*}(\lambda)) = \mathcal{L}_{\text{Valid.}}(w^*(\lambda))$, because $w_{\phi^*}(\lambda) = w^*(\lambda)$. Thus, under mild
 123 conditions, we will learn a best-response in the support of the hyperparameter distribution. If the
 124 best-response is differentiable, then there is a neighborhood about each hyperparameter where the
 125 best-response is approximately linear. If the support of the hyperparameter distribution is this neigh-
 126 borhood, then we can learn the best-response locally with linear regression.

127 In practice, there are no guarantees about the network being a universal approximator or the finite-
 128 time convergence of optimization. The optimal hypernetwork will depend on the hyperparameter
 129 distribution $p(\lambda)$, not just the support of this distribution. We appeal to experimental results that our
 130 method is feasible in practice.

131 2.4 Jointly train parameters and hyperparameters

132 Theorem 2.1 holds for any $p(\lambda)$. In practice, we should
 133 choose a $p(\lambda)$ that puts most of its mass on promising
 134 hyperparameter values, because it may not be possible
 135 to learn a best-response for all hyperparameters due to
 136 limited hypernetwork capacity. Thus, we propose Al-
 137 gorithm 3, which only tries to match a best-response lo-
 138 cally. We introduce a “current” hyperparameter $\hat{\lambda}$, which
 139 is updated each iteration and a conditional hyperparame-
 140 ter distribution, $p(\lambda | \hat{\lambda})$, which only puts mass close to $\hat{\lambda}$.
 141 This allows us to focus on promising hyperparameters to
 142 transition to from our current $\hat{\lambda}$, which an unconditional
 143 $p(\lambda)$ does not, which is critical for allowing the use of a
 144 reduced capacity hypernet, making large problems fea-
 145 sible.

Algorithm 4 Simplified joint training of hypernetwork and hyperparameters

initialize $\phi, \hat{\lambda}$
loop
 $\mathbf{x} \sim \text{Training data}, \mathbf{x}' \sim \text{Validation data}$
 $\phi \leftarrow \alpha \nabla_{\phi} \mathcal{L}_{\text{Train}}(w_{\phi}(\hat{\lambda}), \hat{\lambda}, \mathbf{x})$
 $\hat{\lambda} \leftarrow \beta \nabla_{\hat{\lambda}} \mathcal{L}_{\text{Valid.}}(w_{\phi}(\hat{\lambda}), \mathbf{x}')$
 Return $\hat{\lambda}, w_{\phi}(\hat{\lambda})$

Algorithm 4 builds on Algorithm 3 by using gradient updates on $\hat{\lambda}$ as a source of noise. This variant does not have asymptotic guarantees, but performs similarly to Algorithm 3 in practice.

Algorithm 3 combines the two phases of Algorithm 2. Instead of first learning a hypernetwork that can output weights for any hyperparameter then optimizing the hyperparameters, Algorithm 3 only samples hyperparameters near the current guess. This means the hypernetwork just has to be trained to estimate good enough weights for a small set of hyperparameters. There is an extra cost of having to re-train the hypernetwork each time we update $\hat{\lambda}$. The locally-trained hypernetwork can then be used to provide gradients to update the hyperparameters based on validation set performance.

How simple can we make the hypernetwork, and still obtain useful gradients to optimize hyperparameters? Suppose the hypernetwork is a (reduced-rank) linear function of the hyperparameters and the conditional hyperparameter distribution is $p(\lambda|\hat{\lambda}) = \mathcal{N}(\hat{\lambda}, \sigma\mathbb{I})$ for a small σ . This hypernetwork learns a tangent hyperplane to a best-response function and only needs to make minor adjustments at each step if the hyperparameter updates are sufficiently small. Algorithm 4 uses simultaneous updates and achieves nearly identical results to Algorithm 3 in our experiments.

3 Related work

Our work is complementary to the SMASH algorithm of Brock et al. (2017), with section 2 discussing our differences.

Model-free approaches Model-free approaches use only trial-and-error to explore the hyperparameter space. Simple model-free approaches applied to hyperparameter optimization include grid search and random search (Bergstra & Bengio, 2012). Hyperband (Li et al., 2016) combines bandit approaches with modeling the learning procedure.

Model-based approaches Model-based approaches try to build a surrogate function, which can allow gradient-based optimization or active learning. A common example is Bayesian optimization. Freeze-thaw Bayesian optimization can condition on partially-optimized model performance.

Optimization-based approaches Attempts have been made to directly approximate gradients of the validation loss with respect to hyperparameters. Domke (2012) proposes to differentiate through unrolled optimization to approximate best-responses in nested optimization and Maclaurin et al. (2015a) differentiate through unrolled learning procedures. DrMAD (Fu et al., 2016) approximates differentiating through an unrolled learning procedure to relax memory usage for deep neural networks. HOAG (Pedregosa, 2016) finds hyperparameter gradients with implicit differentiation by deriving an implicit equation for the gradient with optimality conditions. Franceschi et al. (2017) study forward and reverse-mode differentiation for constructing hyperparameter gradients. Also, Feng & Simon (2017) establish conditions where the validation loss of best-responding weights are smooth, allowing gradient-based hyperparameter training. The $T1 - T2$ method of Luketina et al. (2016) provides an algorithm for stochastic gradient-based optimization of hyperparameters. Convergence of $T1 - T2$ depends on approximating the Hessian of the training loss for parameters with the identity matrix.

Game theory Best-response functions are extensively studied as a solution concept in discrete and continuous multi-agent games (e.g., Fudenberg & Levine (1998)). Games where learning a best-response can be applied include adversarial training (Goodfellow et al., 2014), or Stackelberg competitions (e.g., Brückner & Scheffer (2011)). For adversarial training, the analog of our method is a discriminator who observes the generator’s parameters.

4 Experiments

In our experiments, we examine the standard example of stochastic gradient-based optimization of neural networks, with weight regularization penalty. Some gradient-based methods explicitly use the gradient of a loss, while others use the gradient of a learned surrogate loss. Hyper-training learns and substitutes a surrogate best-response function into a real loss. We may contrast our algorithm with methods learning the loss like Bayesian optimization, gradient-based methods only handling hyperparameters that affect the training loss and gradient-based methods which can handle optimization parameters. The best comparison for hyper-training is to gradient-based methods which

only handle parameters affecting the training loss because other methods apply to a more general set of problems. Here, we write the training and validation losses as:

$$\mathcal{L}_{\text{Train}}(w, \lambda) = \mathbb{E}_{\mathbf{x} \sim \text{Train}} \left[\mathcal{L}_{\text{Pred}}(\mathbf{x}, w) \right] + \mathcal{L}_{\text{Reg}}(w, \lambda),$$

$$\mathcal{L}_{\text{Valid.}}(w) = \mathbb{E}_{\mathbf{x} \sim \text{Valid.}} \left[\mathcal{L}_{\text{Pred}}(\mathbf{x}, w) \right]$$

In all experiments, Algorithms 2 or 3 are used to optimize weights with a mean squared error on MNIST (LeCun et al., 1998) with \mathcal{L}_{Reg} as an L_2 weight decay penalty weighted by $\exp(\lambda)$. The elementary model has 7,850 weights. All hidden units in the hypernetwork have a ReLU activation (Nair & Hinton, 2010) unless otherwise specified. Autograd (Maclaurin et al., 2015b) was used to compute all derivatives. For each experiment, the minibatch samples 2 pairs of hyperparameters and up to 1,000 training data points. We used Adam for training the hypernetwork and hyperparameters, with a step size of 0.0001. We ran all experiments on a CPU.

4.1 Learning a global best-response

Our first experiment, shown in Figure 3, demonstrates learning a global approximation to a best-response function using Algorithm 2. To make visualization of the regularization loss easier, we use 10 training data points to exacerbate overfitting. We compare the performance of weights output by the hypernetwork to those trained by standard cross-validation (Algorithm 1). Thus, elementary weights were randomly initialized for each hyperparameter choice and optimized using Adam (Kingma & Ba, 2014) for 1,000 iterations with a step size of 0.0001.

When training the hypernetwork, hyperparameters were sampled from a broad Gaussian distribution: $p(\lambda) = \mathcal{N}(0, 1.5)$. The hypernetwork has 50 hidden units which results in 400,450 parameters of the hypernetwork.

The minimum of the best-response in Figure 3 is near the real minimum of the validation loss, which shows a hypernetwork can satisfactorily approximate a global best-response function in small problems.

4.2 Learning a local best-response

Figure 4 shows the same experiment, but using the Algorithm 3. The fused updates result in finding a best-response approximation whose minimum is the actual minimum faster than the prior experiment. The conditional hyperparameter distribution is given by $p(\lambda|\hat{\lambda}) = \mathcal{N}(\hat{\lambda}, 0.00001)$. The hypernetwork is a linear model, with only 15,700 weights. We use the same optimizer as the global best-response to update both the hypernetwork and the hyperparameters.

Again, the minimum of the best-response at the end of training minimizes the validation loss. This experiment shows that using only a locally-trained linear best-response function can give sufficient gradient information to optimize hyperparameters on a small problem. Algorithm 3 is also less computationally




Figure 3: The validation loss of a neural net, estimated by cross-validation (crosses) or by a hypernetwork (line), which outputs 7,850-dimensional network weights. Cross-validation requires optimizing from scratch each time. The hypernetwork can be used to evaluate the validation loss cheaply.




Figure 4: The losses of a neural net, estimated by cross-validation or a hypernetwork, which outputs 7,850-dimensional network weights. A linear hypernetwork which has limited capacity making it only accurate where the hyperparameter distribution puts mass.

4.3 Hyper-training and unrolled optimization

To compare hyper-training with other gradient-based hyperparameter optimization methods, we train models with 7,850 hyperparameters and a separate L_2 weight decay applied to each weight in a 1 layer (linear) model. The conditional hyperparameter distribution and optimizer for the hypernetwork and hyperparameters is the same the prior experiment. We factorize the weights for the model by selecting a hypernetwork with 10 hidden units. The factorized linear hypernetwork has 10 hidden units giving 164,860 weights. Each hypernetwork iteration is $2 \cdot 10$ times as expensive as an iteration on just the model because there is the same number of hyperparameters as model parameters and we have a 10 unit bottleneck.

Figure 5, top, shows that Algorithm 3 converges more quickly than the unrolled reverse-mode optimization introduced in Maclaurin et al. (2015a) and implemented by Franceschi et al. (2017). Hyper-training reaches sub-optimal solutions because of limitations on how many hyperparameters can be sampled for each update but overfits validation data less than unrolling. Standard Bayesian optimization cannot be scaled to this many hyperparameters. Thus, this experiment shows Algorithm 3 can efficiently partially optimize thousands of hyperparameters. It may be useful to combine these methods by using a hypernetwork to output initial parameters and then unrolling several steps of optimization to differentiate through.

4.4 Optimizing with deeper networks

To see if we can optimize deeper networks with hyper-training we optimize models with 1, 2, and 3 layers and a separate L_2 weight decay applied to each weight. The conditional hyperparameter distribution and optimizer for the hypernetwork and hyperparameters is the same the prior experiment. We factorize the weights for each model by selecting a hypernetwork with 10 hidden units.

Figure 5, bottom, shows that Algorithm 3 can scale to networks with multiple hidden layers and outperform hand-tuned settings. As we add more layers the difference between validation loss and testing loss decreases, and the model performs better on the validation set. Future work should compare other architectures like recurrent or convolutional networks. Additionally, note that more layers perform with lesser training (not shown), validation, and test losses, instead of lower training loss and higher validation or test loss. This performance indicates that using weight decay on each weight could be a prior for generalization, or that hyper-training enforces another useful prior like the continuity of a best-response.

4.5 Estimating weights versus estimating loss

Our approach differs from Bayesian optimization which attempts to directly model the validation loss of optimized weights, where we try to learn to predict optimal weights. In this experiment, we untangle the reason for the better performance of our method: Is it because of a better inductive bias, or because our way can see more hyperparameter settings during optimization?




Figure 5: Validation and test losses during hyperparameter optimization with a separate L_2 weight decay applied to each weight in the model. Thus, models with more parameters have more hyperparameters. *Top*: We solve the 7,850-dimensional hyperparameter optimization problem with a linear model and multiple algorithms. Hypernetwork-based optimization converges to a sub-optimal solution faster than unrolled optimization from Maclaurin et al. (2015a). *Bottom*: Hyper-training is applied different layer configurations in the model.

297

298 First, we constructed a hyper-training set:
 299 We optimized 25 sets of weights to completion, given randomly-sampled hyperparameters. We chose 25 samples since that
 300 is the regime in which we expect Gaussian process-based approaches to have the most significant advantage. We also constructed a validation set of 10,215 (optimized weight, hyperparameter) tuples generated in the same manner. **We then**
 302 fit a Gaussian process (GP) regression model with an RBF kernel from sklearn on the validation loss data. **A hypernetwork**
 303 **is fit to the same set of hyperparameters and data.** Finally, we optimize another hypernetwork using Algorithm 2, for
 304 the same amount of time as building the GP training set. The two hypernetworks were linear models and trained with the same optimizer parameters as the 7,850-dimensional hyperparameter optimization.

319 Figure 6 shows the distribution of prediction errors of these three models. We can see that the Gaussian process tends to underestimate loss. The hypernetwork trained with the same small fixed set of examples tends to overestimate loss. We conjecture that this is due to the hypernetwork producing bad weights in regions where it doesn't have enough training data. Because the hypernetwork must provide actual weights to predict the validation loss, poorly-fit regions will overestimate the validation loss. Finally, the hypernetwork trained with Algorithm 2 produces errors tightly centered around 0. The main takeaway from this experiment is a hypernetwork can learn more accurate surrogate functions than a GP for equal compute budgets because it views (noisy) evaluations of more points.




Figure 6: Comparing three approaches to inferring validation loss. *First column:* A Gaussian process, fit on 25 hyperparameters and the corresponding validation losses. *Second column:* A hypernetwork, fit on the same 25 hyperparameters and the corresponding optimized weights. *Third column:* Our proposed method, a hypernetwork trained with stochastically sampled hyperparameters. *Top row:* The distribution of inferred and true losses. The diagonal black line is where predicted loss equals true loss. *Bottom row:* The distribution of differences between inferred and true losses. The Gaussian process often under-predicts the true loss, while the hypernetwork trained on the same data tends to over-predict the true loss.

333 5 Conclusions and Future Work

334 In this paper, we addressed the question of tuning hyperparameters using gradient-based optimization, by replacing the training optimization loop with a differentiable hypernetwork. We gave a theoretical justification that sufficiently large networks will learn the best-response for all hyperparameters viewed in training. We also presented a simpler and more scalable method that jointly optimizes
 336 both hyperparameters and hypernetwork weights, allowing our method to work with manageably-sized hypernetworks.

340 Experimentally, we showed that hypernetworks could provide a better inductive bias for hyperparameter optimization than Gaussian processes fitting the validation loss empirically.

342 There are many directions to extend the proposed methods. For instance, the hypernetwork could be composed with several iterations of optimization, as an easily-differentiable fine-tuning step. Or, hypernetworks could be incorporated into meta-learning schemes, such as MAML (Finn et al., 2017), which finds weights that perform a variety of tasks after unrolling gradient descent.

346 We also note that the prospect of optimizing thousands of hyperparameters raises the question of hyper-regularization, or regularization of hyperparameters.

References

- Bergstra, James and Bengio, Yoshua. Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13(Feb):281–305, 2012.
- Blundell, Charles, Cornebise, Julien, Kavukcuoglu, Koray, and Wierstra, Daan. Weight uncertainty in neural networks. *arXiv preprint arXiv:1505.05424*, 2015.
- Brock, Andrew, Lim, Theodore, Ritchie, JM, and Weston, Nick. Smash: One-shot model architecture search through hypernetworks. *arXiv preprint arXiv:1708.05344*, 2017.
- Brückner, Michael and Scheffer, Tobias. Stackelberg games for adversarial prediction problems. In *Proceedings of the 17th ACM SIGKDD international conference on Knowledge discovery and data mining*, pp. 547–555. ACM, 2011.
- Domke, Justin. Generic methods for optimization-based modeling. In *Artificial Intelligence and Statistics*, pp. 318–326, 2012.
- Feng, Jean and Simon, Noah. Gradient-based regularization parameter selection for problems with non-smooth penalty functions. *arXiv preprint arXiv:1703.09813*, 2017.
- Finn, Chelsea, Abbeel, Pieter, and Levine, Sergey. Model-agnostic meta-learning for fast adaptation of deep networks. *arXiv preprint arXiv:1703.03400*, 2017.
- Franceschi, Luca, Donini, Michele, Frasconi, Paolo, and Pontil, Massimiliano. Forward and reverse gradient-based hyperparameter optimization. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pp. 1165–1173. PMLR, 2017.
- Fu, Jie, Luo, Hongyin, Feng, Jiashi, Low, Kian Hsiang, and Chua, Tat-Seng. Drmad: distilling reverse-mode automatic differentiation for optimizing hyperparameters of deep neural networks. *arXiv preprint arXiv:1601.00917*, 2016.
- Fudenberg, Drew and Levine, David K. *The theory of learning in games*, volume 2. MIT press, 1998.
- Goodfellow, Ian, Pouget-Abadie, Jean, Mirza, Mehdi, Xu, Bing, Warde-Farley, David, Ozair, Sherjil, Courville, Aaron, and Bengio, Yoshua. Generative adversarial nets. In *Advances in neural information processing systems*, pp. 2672–2680, 2014.
- Ha, David, Dai, Andrew, and Le, Quoc V. Hypernetworks. *arXiv preprint arXiv:1609.09106*, 2016.
- Hornik, Kurt. Approximation capabilities of multilayer feedforward networks. *Neural networks*, 4(2):251–257, 1991.
- Kingma, Diederik and Ba, Jimmy. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- LeCun, Yann, Bottou, Léon, Bengio, Yoshua, and Haffner, Patrick. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- Li, Lisha, Jamieson, Kevin, DeSalvo, Giulia, Rostamizadeh, Afshin, and Talwalkar, Ameet. Hyperband: A novel bandit-based approach to hyperparameter optimization. *arXiv preprint arXiv:1603.06560*, 2016.
- Lizotte, Daniel James. *Practical bayesian optimization*. University of Alberta, 2008.
- Luketina, Jelena, Berglund, Mathias, Greff, Klaus, and Raiko, Tapani. Scalable gradient-based tuning of continuous regularization hyperparameters. In *International Conference on Machine Learning*, pp. 2952–2960, 2016.
- Maclaurin, Dougal, Duvenaud, David, and Adams, Ryan. Gradient-based hyperparameter optimization through reversible learning. In *International Conference on Machine Learning*, pp. 2113–2122, 2015a.

393 Maclaurin, Dougal, Duvenaud, David, and Adams, Ryan P. Autograd: Effortless gradients in numpy.
394 In *ICML 2015 AutoML Workshop*, 2015b.

395 Nair, Vinod and Hinton, Geoffrey E. Rectified linear units improve restricted boltzmann machines.
396 In *Proceedings of the 27th international conference on machine learning (ICML-10)*, pp. 807–
397 814, 2010.

398 Pedregosa, Fabian. Hyperparameter optimization with approximate gradient. In *International Con-*
399 *ference on Machine Learning*, pp. 737–746, 2016.

400 Rasmussen, Carl Edward and Williams, Christopher KI. *Gaussian processes for machine learning*,
401 volume 1. MIT press Cambridge, 2006.

402 Snoek, Jasper, Larochelle, Hugo, and Adams, Ryan P. Practical bayesian optimization of machine
403 learning algorithms. In *Advances in neural information processing systems*, pp. 2951–2959, 2012.

404 Swersky, Kevin, Snoek, Jasper, and Adams, Ryan Prescott. Freeze-thaw bayesian optimization.
405 *arXiv preprint arXiv:1406.3896*, 2014.